


# Deep Learning Topics: Special Connections, Encoder-Decoders, Transformers

Paul Rodriguez, PhD  
(SDSC)  
August 2024



# Outline

- **Part I**

**Gate connection idea**

**Skip and Residual connections**

**Programing connections and Keras Model API**

**Encoder-Decoder (Autoencoder)**

**Exercise MNIST Autoencoder**




**Autoencoder with Stable Diffusion**

- **Part II**

**Attention Head and Transformers**



# Recall: the logistic unit

$$f(x) = \frac{1}{1 + \exp(-(b + wx))}$$



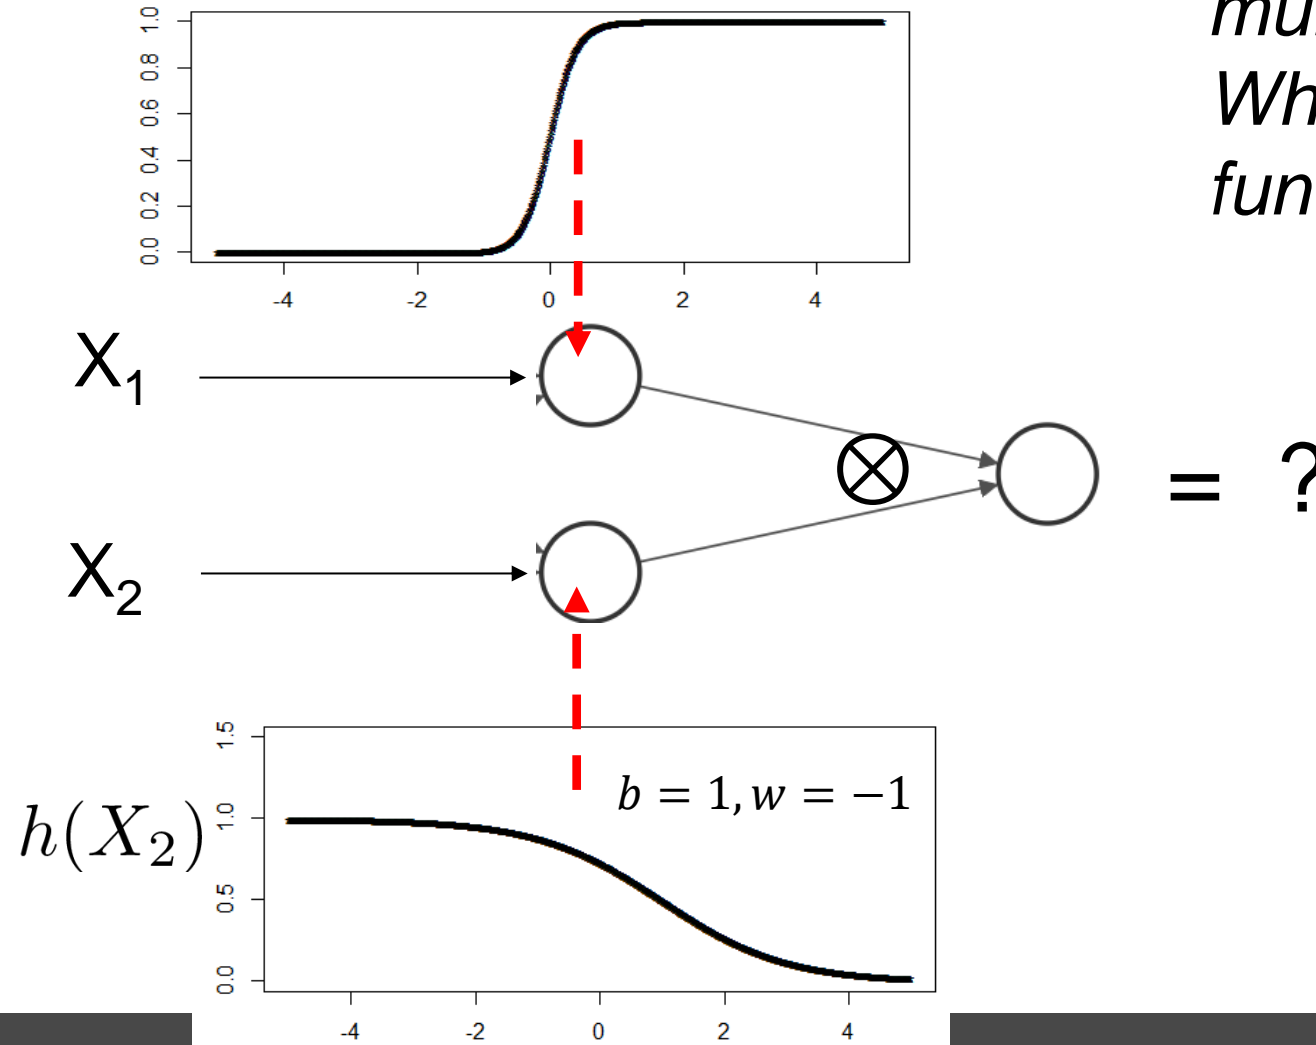
# Example: 1 input into 2 logistic units with these activations

*If you add these 2 units into a final output unit what would the output function look like?*



# Example: 2 input into 2 logistic units with these activations


*What if you multiply these?  
What is the output function doing?*



# Example: 2 input into 2 logistic units with these activations

*What if you multiply these?*


*For linear activation, what is the output function doing ?*



$$= \begin{cases} 0 & \text{if } X_1 < 0 \\ h(X_2) & \text{if } X_1 > 0 \end{cases}$$


$X_1$  "gates"  $X_2$  activation

# Example: 2 pairs of gates together



Let  $h_1()$  be logistic function to be learned.


# Example: 2 pairs of gates together



Let  $h_1()$  be logistic function to be learned.  
If we set  $h_2(G_2) = (1 - h_1(G_1))$   
then output is probabilistically weighted.




# Example: 2 pairs of gates together



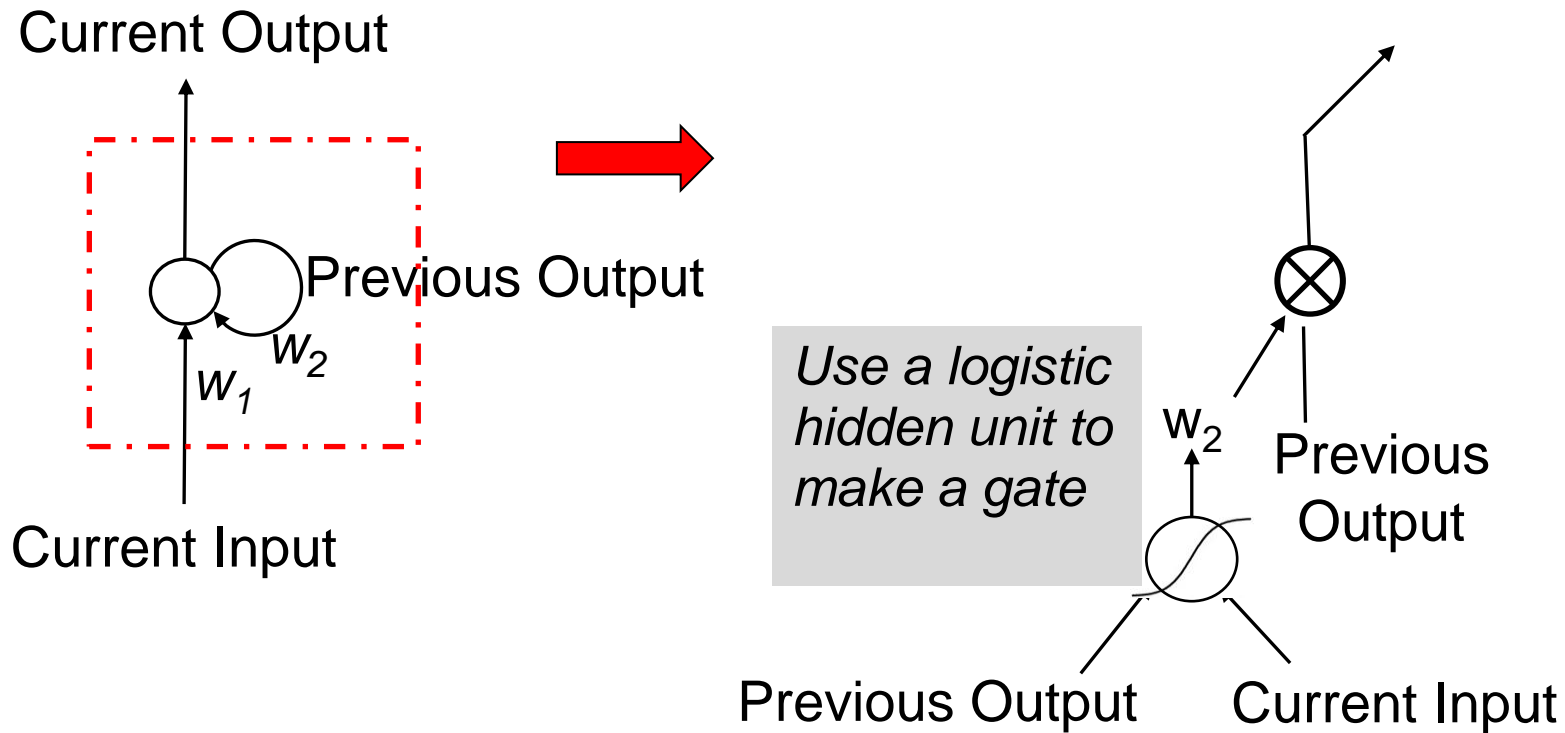
Let  $h_1()$  be logistic function to be learned.  
If we set  $h_2(G_2) = (1 - h_1(G_1))$   
then output is probabilistically weighted.

***Q: Where should the  $G_1$  value come from?***

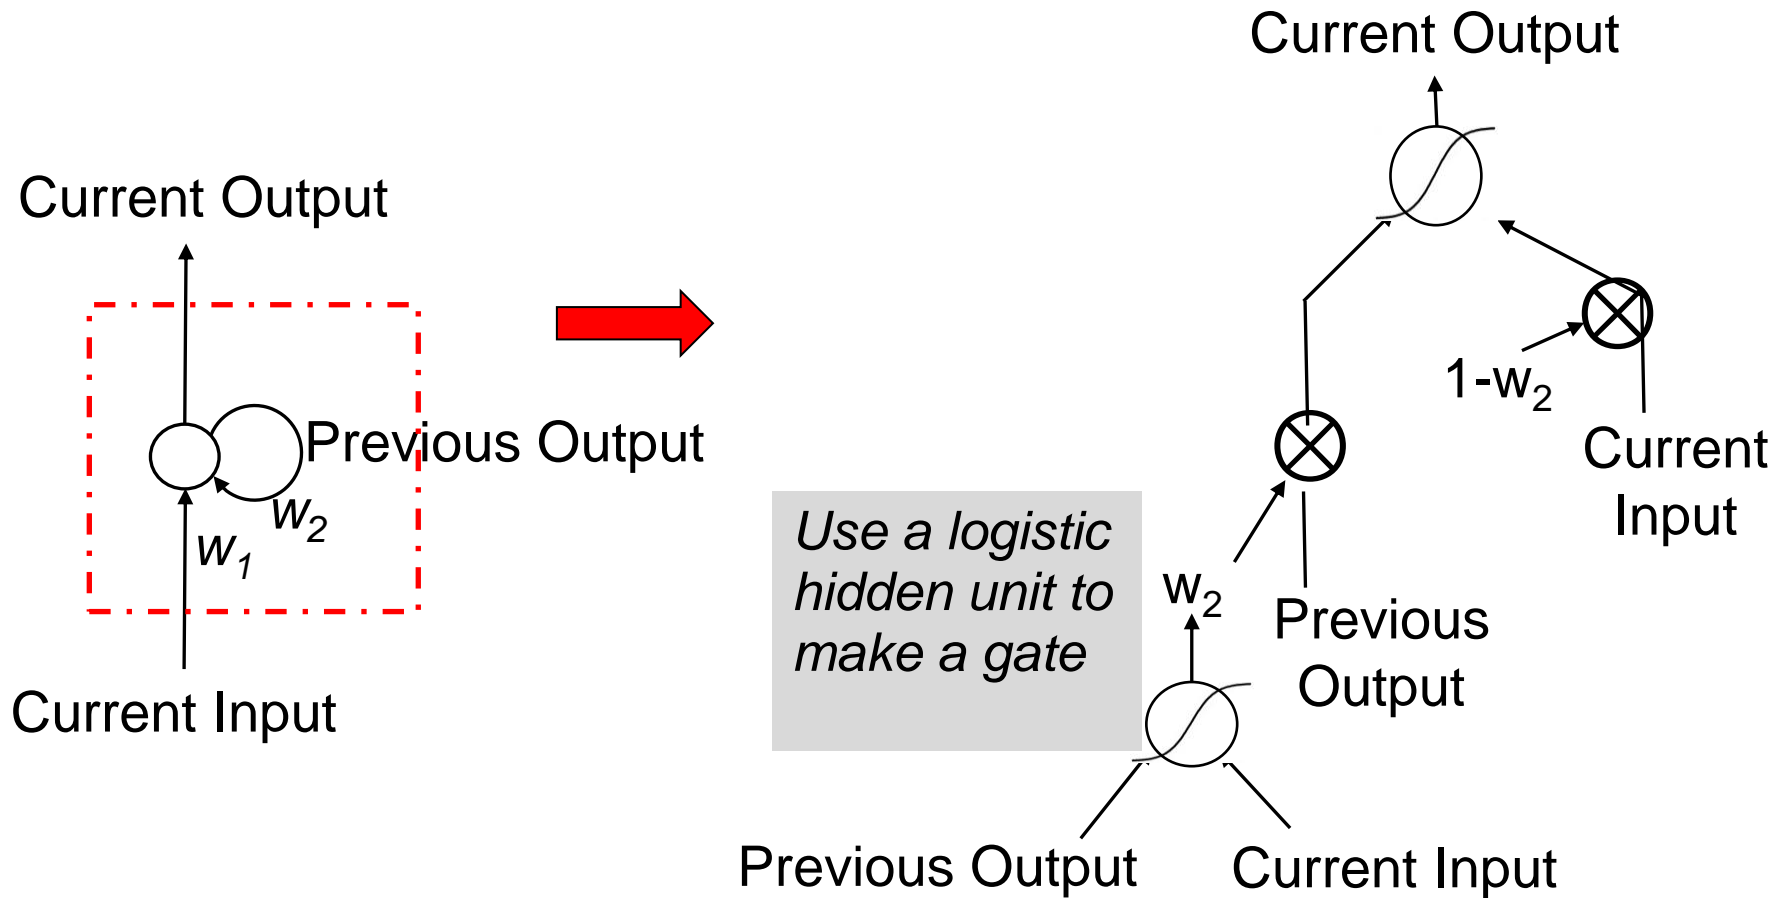
# A recurrent unit for sequence learning can be replaced by a gated unit



# A recurrent unit for sequence learning can be replaced by a gated unit



# A recurrent unit for sequence learning can be replaced by a gated unit




*Use the gate to either keep previous output or update it with current input*


*Use a logistic hidden unit to make a gate*

*'Gated Recurrent Unit'  
Cho, Bengio 2015*

# Redrawing the gate for sets of hidden units



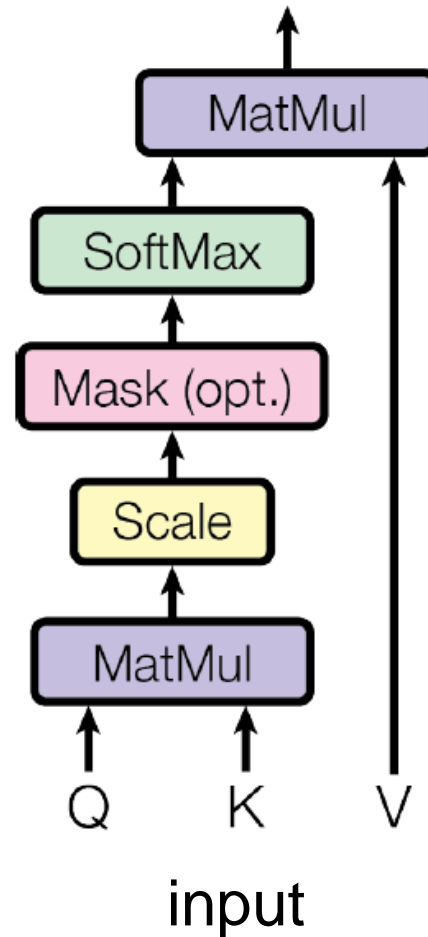
# Use softmax for $G(X)$ to get gating weights



**Recall: softmax  
normalizes outputs  
into probability  
weights**

## Scaled Dot-Product Attention (very rough summary)

“Attention” mechanism in language transformers use a softmax gate




The gate is applied to possible Values ( $V$ ) for decoding

Vaswani, et al. 2017  
Attention Is All You Need (for Transformers)

# Skip Connections:

## Recall the Multilayer Perceptron (MLP)


Input layer    Hidden layer    Output layer






# To help the MLP learn directly from input carry input forward


Input layer   Hidden layer   Output layer



# Concatenate input with hidden units into new layer



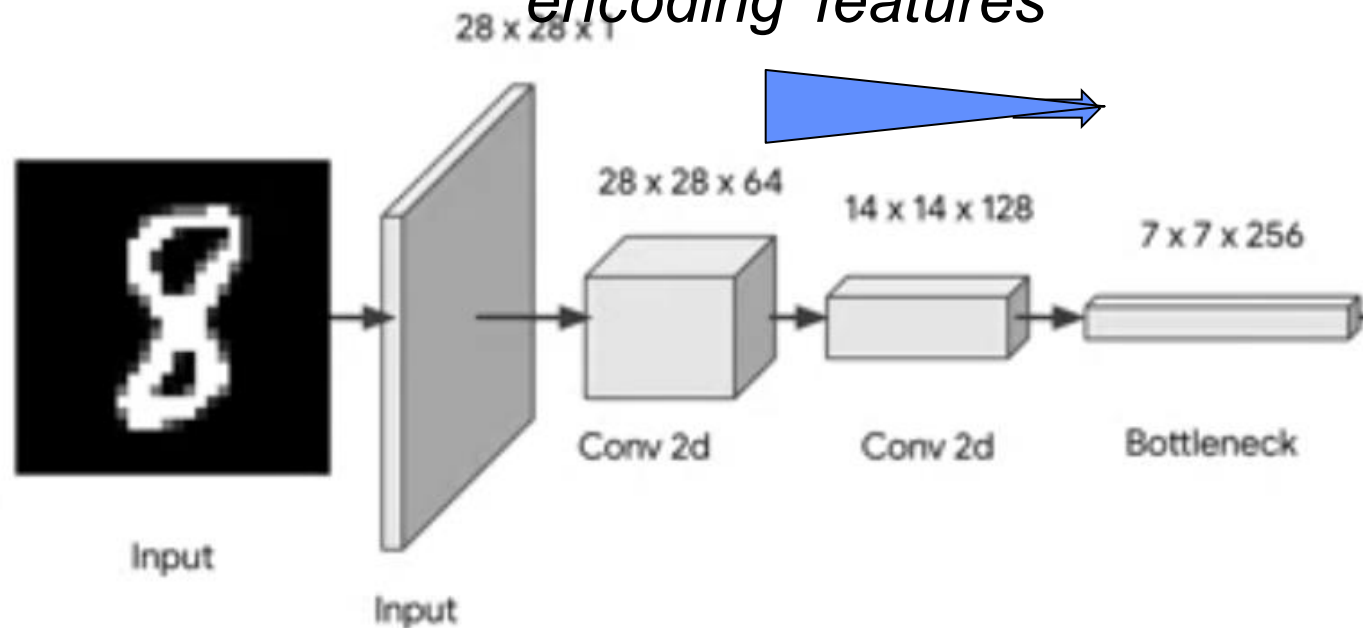
# Can be done for any (or all) previous layer and *skip* any number of layers



# Recall: CNN architecture for MNIST classification

ENCODER

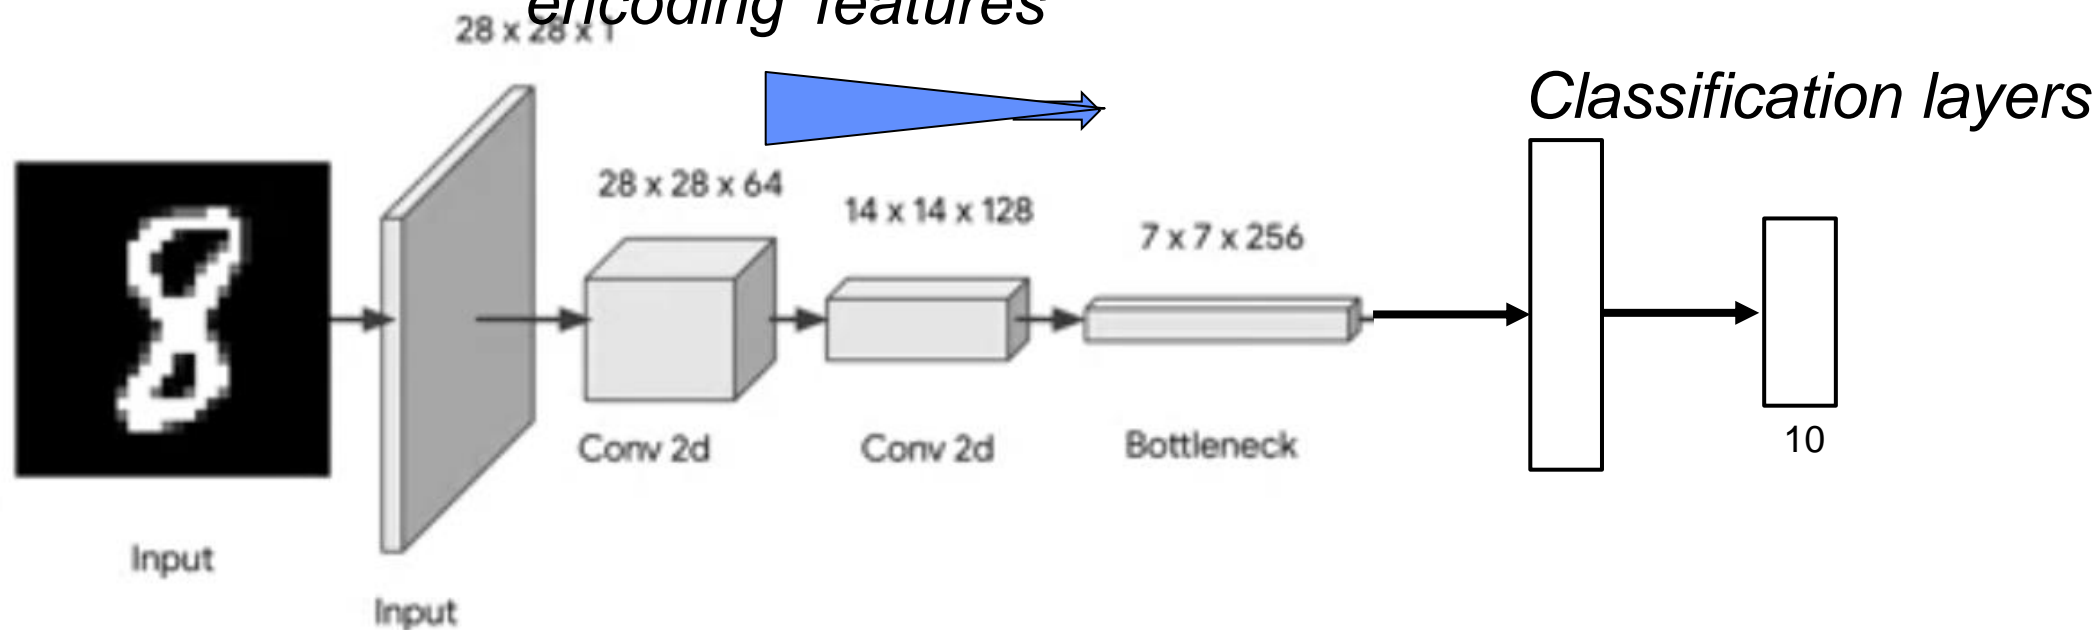
*more feature maps & downsampling :  
'encoding' features*



# Consider: CNN architecture for MNIST classification

ENCODER

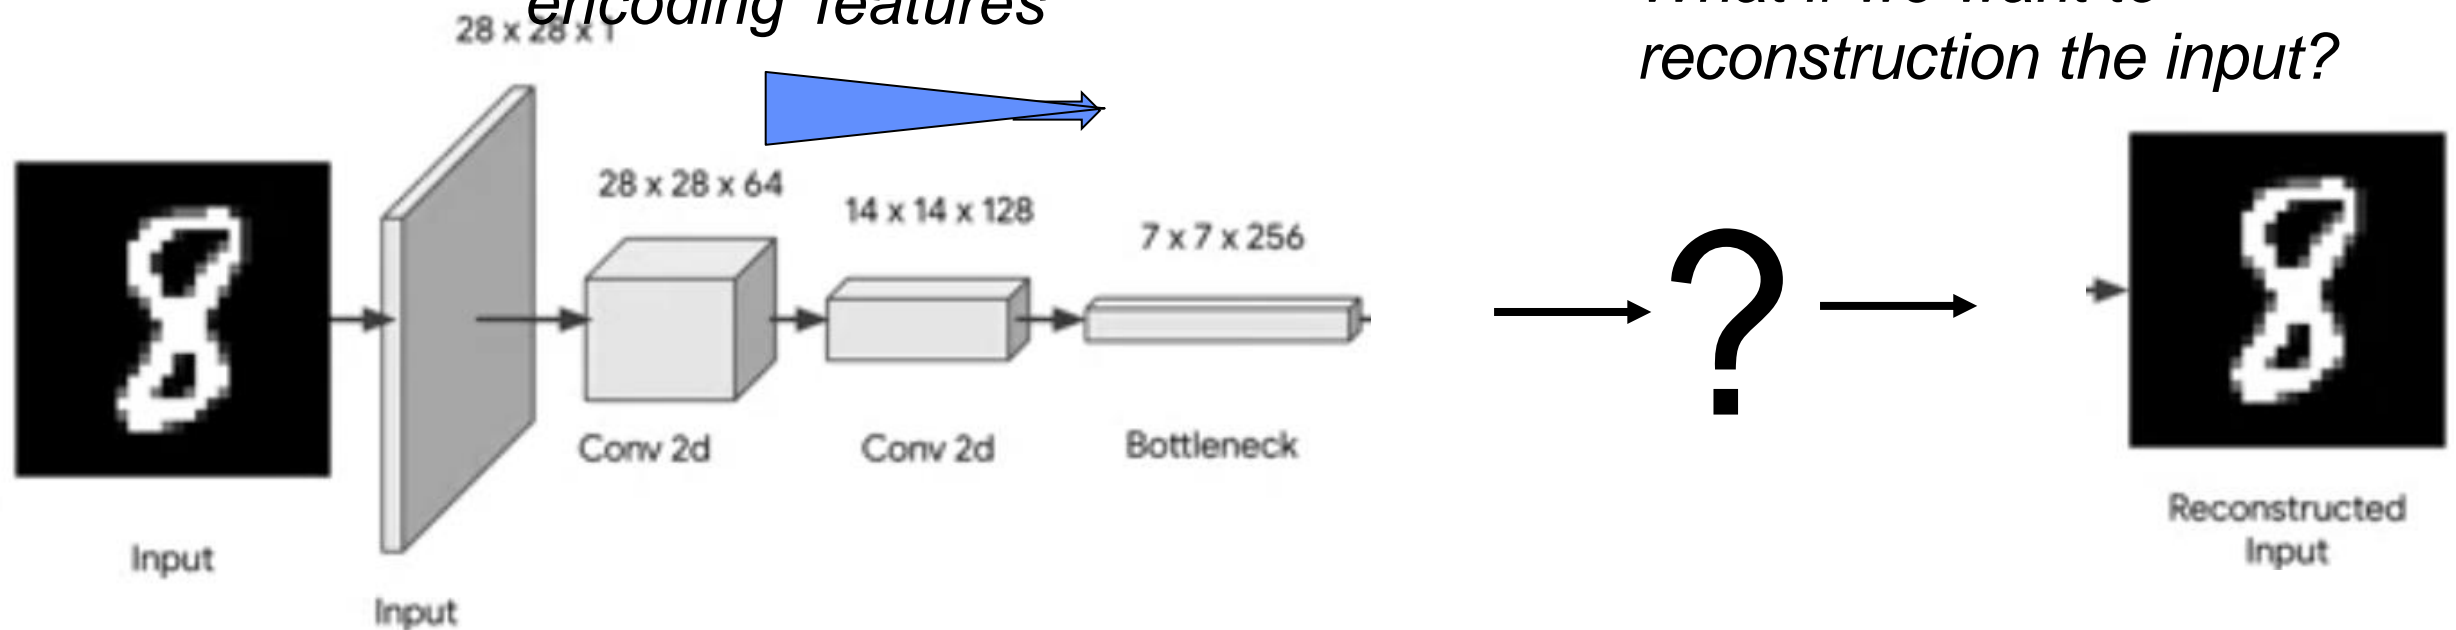
*more feature maps & downsampling :  
'encoding' features*



# A CNN architecture for MNIST autoencoding

ENCODER

*more feature maps & downsampling :  
'encoding' features*




# A CNN architecture for MNIST autoencoding

ENCODER

DECODER

*more feature maps & downsampling :  
'encoding' features*

*What if we want to  
reconstruction the input?*



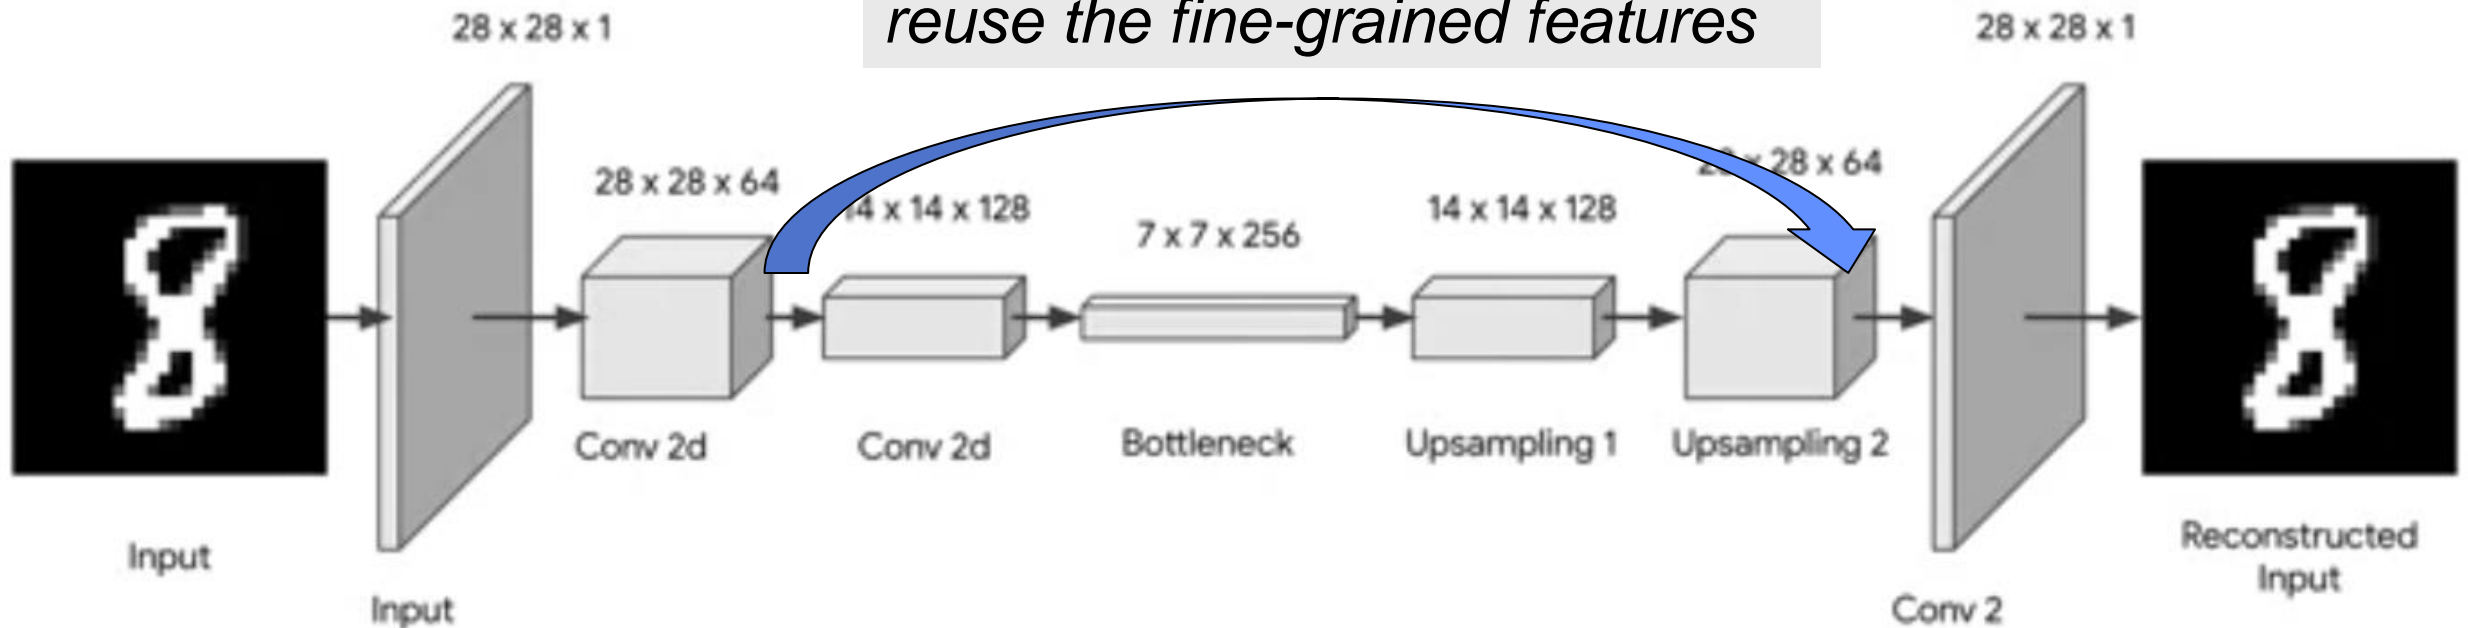
*Use 'upsampling' to reverse  
the encoding – ie. decoding*

# A CNN architecture for MNIST autoencoding

ENCODER

DECODER

*Adding skip connections helps reuse the fine-grained features*



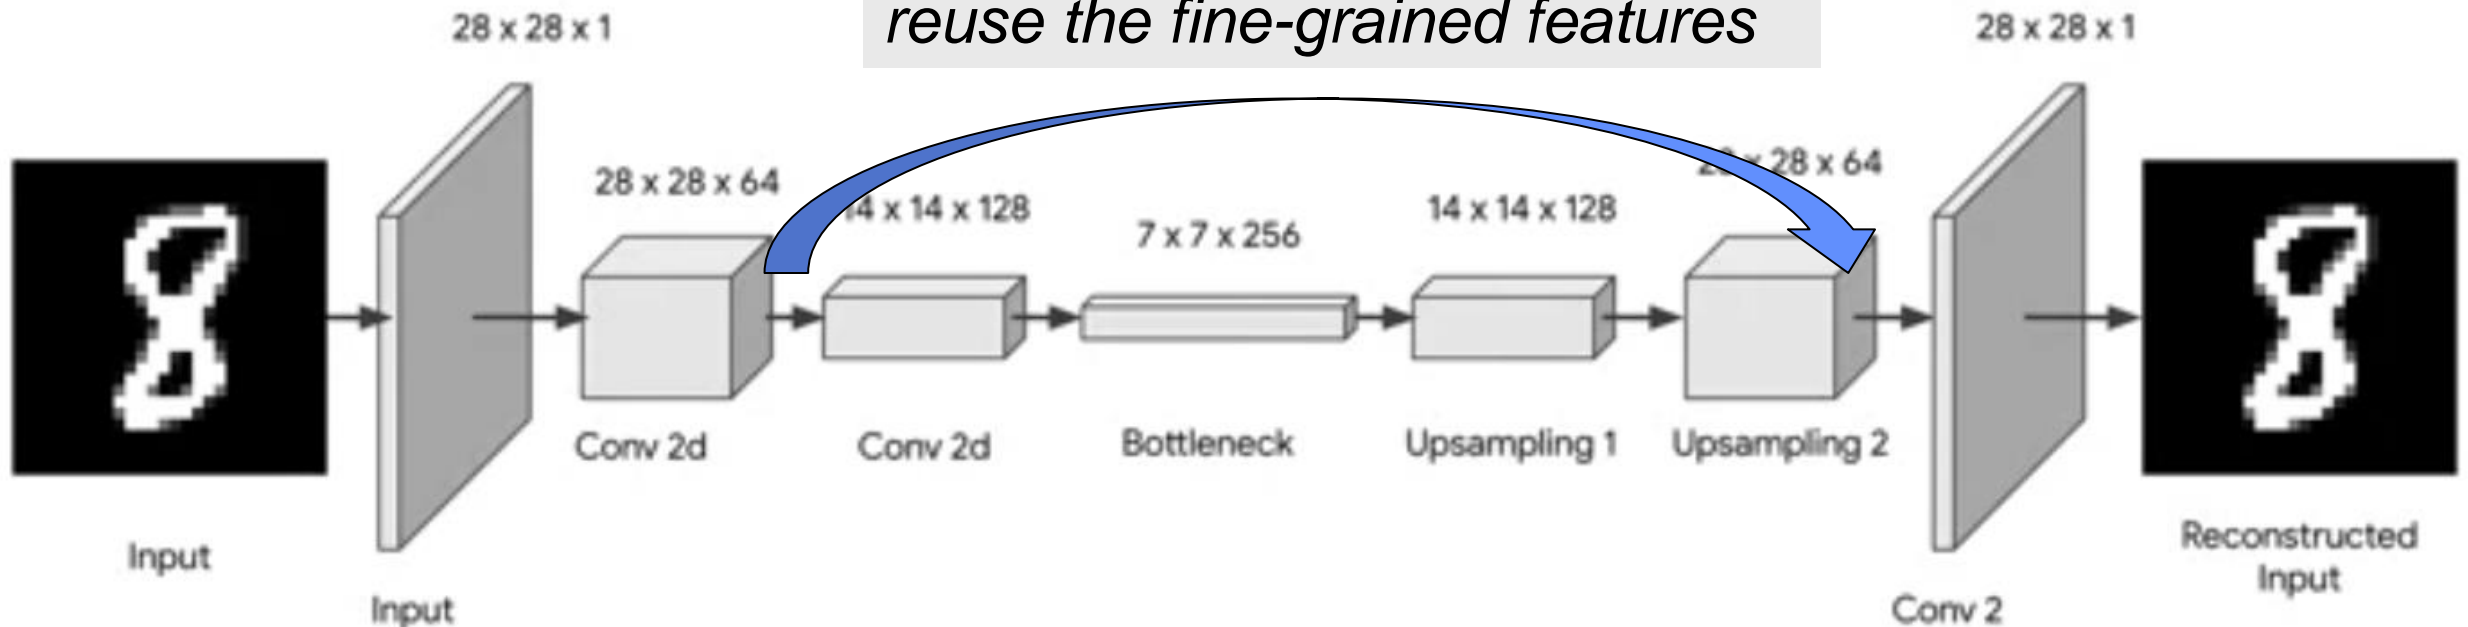


# A CNN architecture for MNIST autoencoding

ENCODER


DECODER

*Adding skip connections helps reuse the fine-grained features*




*NOTE the 28x28x64 encoded maps have to be skipped ahead to where the 28x28x64 decoding maps are – which axis is concatenated?*

# Image Encoder-Decoder is a “UNET” architecture




- **pause**

# Consider: Can we keep adding deep layers?




Given some deep network,  
should I add another layer?  
What should a new layer learn?

# Consider: Can we keep adding deep layers?




If  $H1(X)$  is good then this new layer could be unnecessary,  
Eg  $H2(X)$  should be just  $H1(X)$

# Skip with addition makes a 'residual' connection



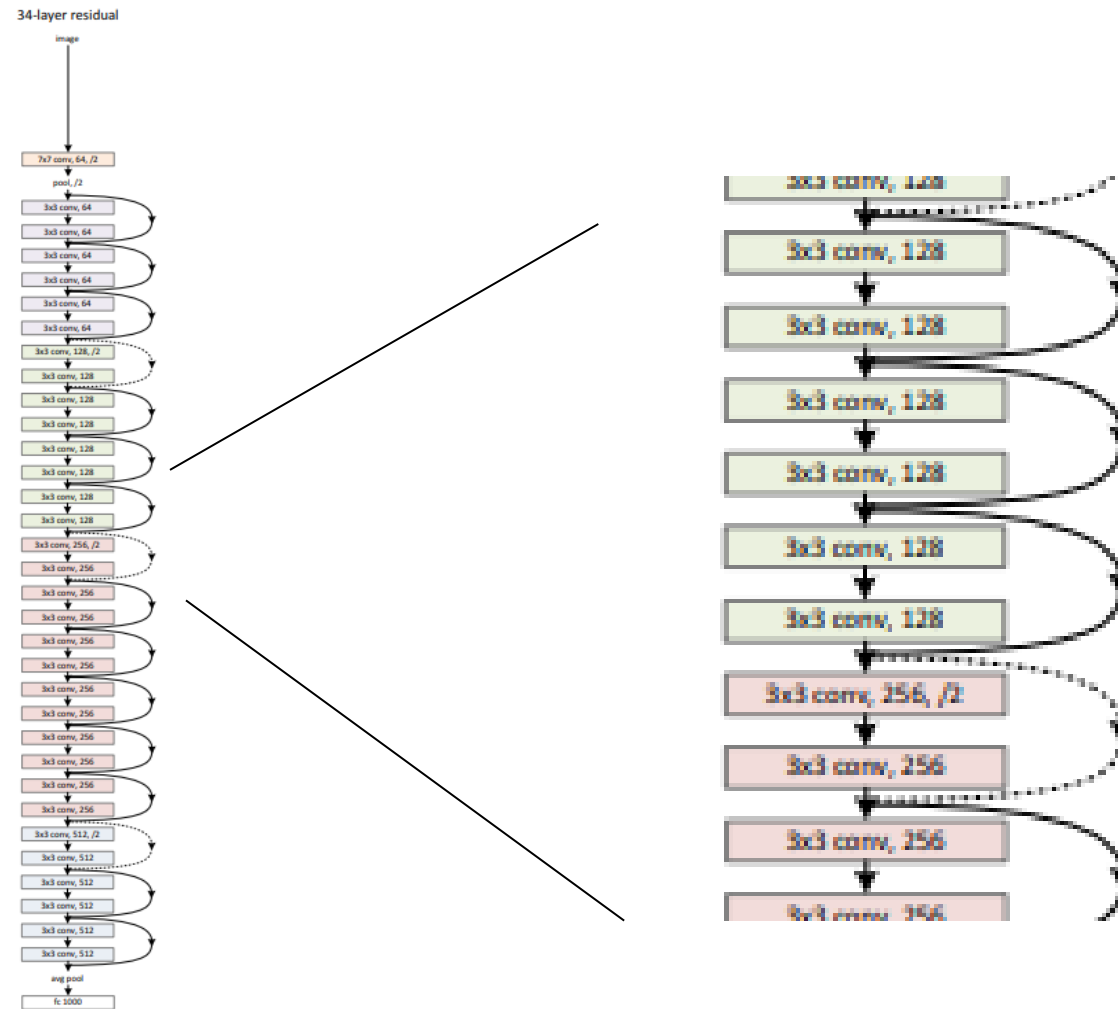
Make it easy for next layer to learn nothing –

# Skip with addition makes a 'residual' connection



Make it easy for next layer to learn nothing –  
e.g. use  $F(X)=H2(X)+H1(X)$  so that  $H2(X)=F(X)-H1(X)$  .  
The  $H2()$  function learned is a residual function

# “Resnet” residual connections help deeper learning




*Deep Residual Learning,  
He et.al, 2015*




# Summary: useful connections for architectures, and the intuitions

Softmax for gating




Recurrent nets,  
attention network

Skip connections  
for feature reuse



UNET, also  
feedforward nets..

Residual connections  
help deeper learning



Resnet, large image  
classification

- **Programing connections and Keras Model API**

# Keras: Sequential API VS Functional API

```
#specify the neural network model and Learning parameters  
my_model = tf.keras.models.Sequential([  
    tf.keras.layers.Flatten(input_shape=(28, 28)),  
    tf.keras.layers.Dense(32,activation='relu'),  
    tf.keras.layers.Dense(10,activation='softmax')])  
my_model.summary()
```

*A sequence of layers:  
the inputs are assumed  
to be in order*

# Keras: Sequential API VS Functional API

```
#specify the neural network model and learning parameters
my_model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(32,activation='relu'),
    tf.keras.layers.Dense(10,activation='softmax')])
my_model.summary()
```

*A sequence of layers:  
the inputs are assumed  
to be in order*

```
#specify the neural network model and learning parameters
inputs = tf.keras.layers.Input(shape=(28, 28, 1,))
inputs_flattened = tf.keras.layers.Flatten()(inputs)
hidden_layer = tf.keras.layers.Dense(32,activation='relu')(inputs_flattened)
output_layer = tf.keras.layers.Dense(10,activation='softmax')(hidden_layer)
```

*A sequence of functions:  
Input layer(s) are specified*

# Keras: Sequential API VS Functional API

```
#specify the neural network model and learning parameters
my_model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(32,activation='relu'),
    tf.keras.layers.Dense(10,activation='softmax')])
my_model.summary()
```

*A sequence of layers:  
the inputs are assumed  
to be in order*

```
#specify the neural network model and learning parameters
inputs = tf.keras.layers.Input(shape=(28, 28, 1,))
inputs_flattened = tf.keras.layers.Flatten()(inputs)
hidden_layer = tf.keras.layers.Dense(32,activation='relu')(inputs_flattened)
output_layer = tf.keras.layers.Dense(10,activation='softmax')(hidden_layer)

my_model = tf.keras.Model(inputs,output_layer)
my_model.summary()
```

*A sequence of functions:  
Input layer(s) are specified*

*The Model() function figures out the full  
path(s) to connect the input(s) to output(s)*

# Keras: Functional API

*A sequence of functions:  
Input layer(s) are specified*

*#specify the neural network model and learning parameters*

```
inputs = tf.keras.layers.Input(shape=(28, 28, 1,))
inputs_flattened = tf.keras.layers.Flatten()(inputs)
hidden_layer = tf.keras.layers.Dense(32, activation='relu')(inputs_flattened)
output_layer = tf.keras.layers.Dense(10, activation='softmax')(hidden_layer)
```

```
my_model = tf.keras.Model(inputs=inputs, outputs=[output_layer, hidden_layer])
```

```
my_model.compile(optimizer=tf.keras.optimizers.Adam(),
                 loss={'output_layer': 'binary_crossentropy',
                      'hidden_layer': 'mae'},
                 ...)
```

*The Model() function can  
also have multiple outputs  
with corresponding loss  
functions.*

# Keras: Functional API

*A sequence of functions:  
Input layer(s) are specified*

*#specify the neural network model and learning parameters*

```
inputs          = tf.keras.layers.Input(shape=(28, 28, 1,))
inputs_flattened = tf.keras.layers.Flatten()(inputs)
hidden_layer    = tf.keras.layers.Dense(32,activation='relu')(inputs_flattened)
output_layer    = tf.keras.layers.Dense(10,activation='softmax')(hidden_layer)
```

```
my_model = tf.keras.Model(inputs =inputs, outputs=[output_layer, hidden_layer])
```

```
my_model.compile(optimizer=tf.keras.optimizers.Adam(),
                  loss={'output_layer': 'binary_crossentropy',
                        'hidden_layer': 'mae'},
                  ...)
```

*The Model() function can  
also have multiple outputs  
with corresponding loss  
functions.*

*This could 'inform' the network to learn some property or constraint*

# Exercise

- MNIST autoencoder, reconstruct digits from noisy inputs
- Add skip connections with concatenation

*Note: make sure outputs from encoding layers are matched up to inputs for decoding layers!*

*i.e. 14x14 encoding feature maps should be concatenated with 14x14 decoding maps*

- Review outputs to see improvements



Login to expanse and start a notebook on  
gpu-shared queue

\$ jupyter-gpu-shared-tensorflow

In jupyter notebook session open the  
MNIST\_Autoencoder notebook

Follow instructions in the notebook

Expanse Dashboard - Expanse Portal CIML22\_DL/

oxford-liftoff-munchkin.expanse-user-content.sdsc.edu/tree/CIML22\_DL

jupyter

Files Running Clusters

Select items to perform actions on them.

0 / CIML22\_DL

..

☐ CIML22\_DL\_MNIST\_INTRO\_v1.ipynb

☐ CIML22\_DL\_MNIST\_INTRO\_v1soltn.ipynb

☐ CIML22\_DL\_MNIST\_wHyperTuner\_v1.ipynb

☐ CIML22\_DL\_MNIST\_wHyperTuner\_v1soltn.ipynb

☒ CIML22\_MNIST\_Autoencoder\_v2.ipynb

☐ CIML22\_MNIST\_Autoencoder\_v2soltn.ipynb

☐ CIML22\_Simple\_MNIST\_SeqandFunc\_v1.ipynb

https://oxford-liftoff-munchkin.expanse-user-content.sdsc.edu/notebooks/CIML22\_DL/CIML22\_DL\_MNIST\_INTRO\_

# Quick overview of code

```
def encoder(inputs):  
    '''Defines the encoder with two Conv2D and max pooling layers.'''  
    conv_1 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), activation='relu', padding='same')(inputs)  
    #padding same produces same output size  
    max_pool_1 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_1) #max pooling does the downsampling
```

Encoder function -

has convolution and pooling layers

```
def decoder(inputs, enc conv1, enc conv2):  
    '''Defines the decoder path to upsample back to the original image size.'''  
    #Notice that padding = same keeps the output same size as input  
  
    conv_1 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), activation='relu', padding='same')(inputs)  
    up_sample_1 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_1)
```

Decoder function -

has up sampling (deconvolution) layers

```
def encoder(inputs):
    """Defines the encoder with two Conv2D and max pooling layers."""
    conv_1 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), activation='relu', padding='same')(inputs)
    #padding same produces same output size
    max_pool_1 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_1) #max pooling does the downsampling

    conv_2 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), activation='relu', padding='same')(max_pool_1)
    max_pool_2 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_2)

    return max_pool_2, conv_1, conv_2
```

Encoder returns final and intermediate layer outputs to be skipped ahead

```
def decoder(inputs, enc conv1, enc conv2):
    """Defines the decoder path to upsample back to the original image size."""
    #Notice that padding = same keeps the output same size as input

    conv_1 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), activation='relu', padding='same')(inputs)
    up_sample_1 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_1)
```

```
def encoder(inputs):
    """Defines the encoder with two Conv2D and max pooling layers."""
    conv_1 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), activation='relu', padding='same')(inputs)
    #padding same produces same output size
    max_pool_1 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_1) #max pooling does the downsampling

    conv_2 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), activation='relu', padding='same')(max_pool_1)
    max_pool_2 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_2)

    return max_pool_2, conv_1, conv_2
```

Encoder returns final and intermediate layer outputs to be skipped ahead

```
def decoder(inputs, enc_conv1, enc_conv2):
    """Defines the decoder path to upsample back to the original image size."""
    #Notice that padding = same keeps the output same size as input

    conv_1 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), activation='relu', padding='s
    up_sample_1 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_1)
```

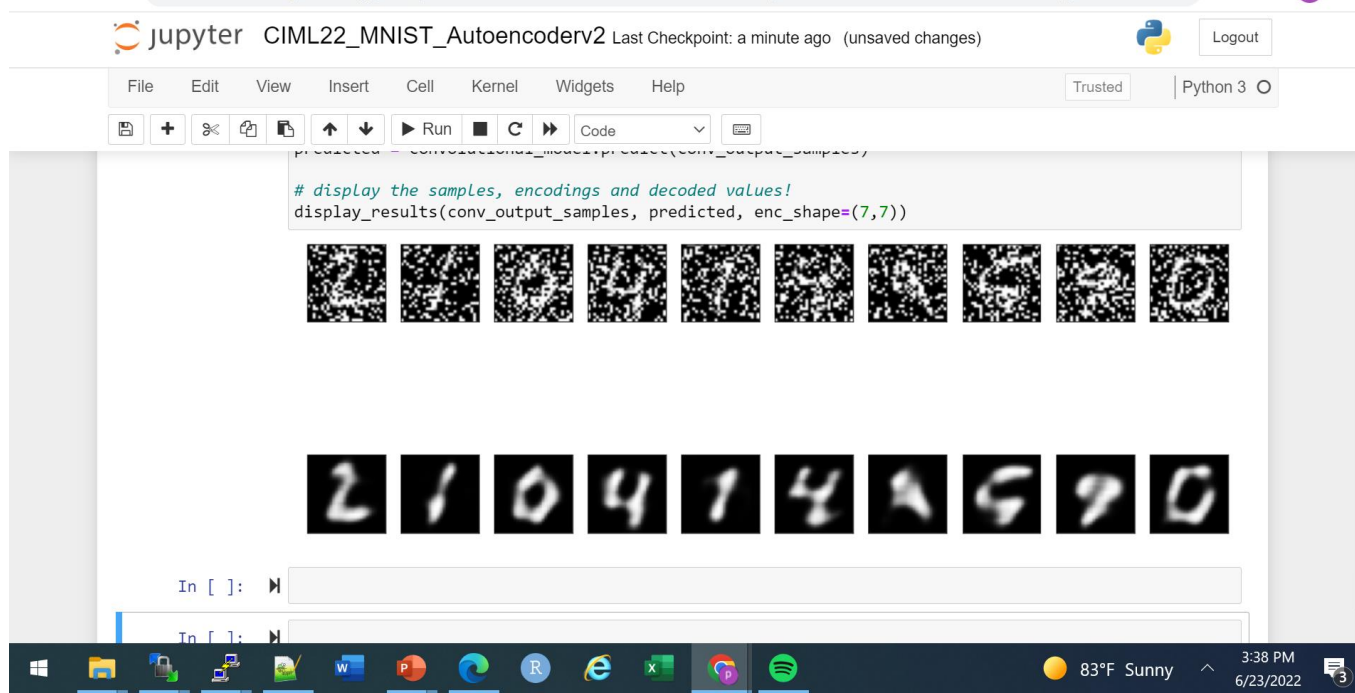
You can pass intermediate layers to decoder,

...

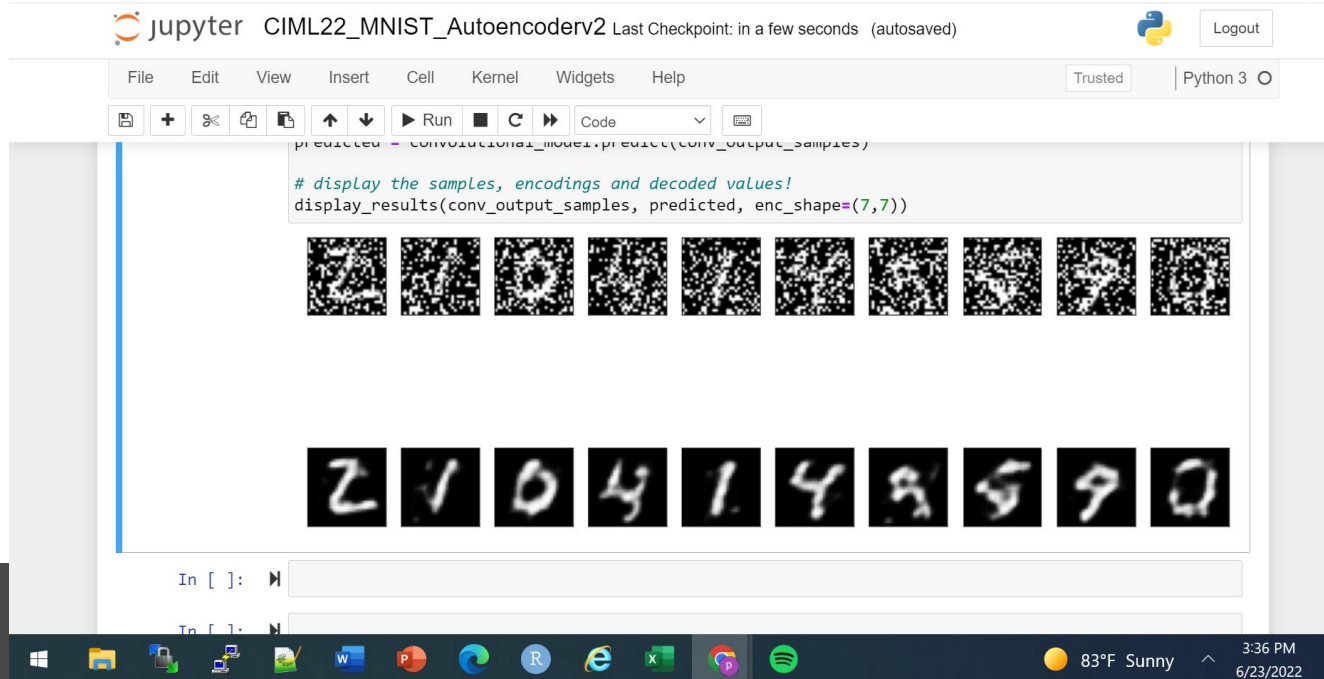
...

then use it in concatenation layer

```
skip_concat_1 = tf.keras.layers.concatenate([up_sample_1, enc_conv2])
conv_2 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), activation='relu', padding='s
# ----->>> and change the input into conv_2
```



With out skip  
20 epochs  
Loss 0.1664



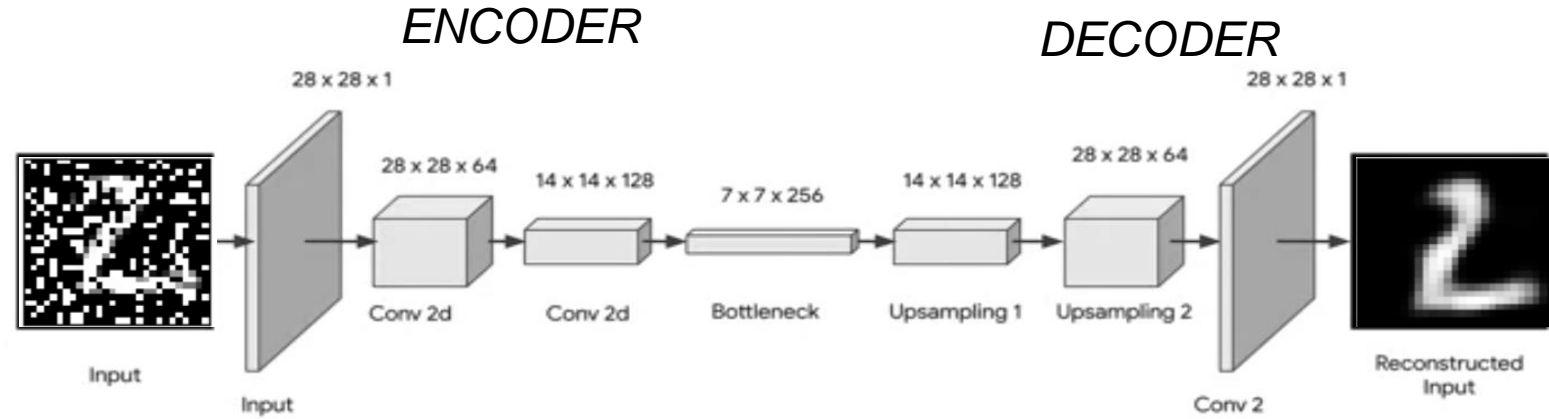
With skip,  
20 epochs loss 0.14

Are the numbers a little bit  
more reconstructed?


# Autoencoding with Stable Diffusion

- Let's introduce the concepts and intuition behind stable diffusion

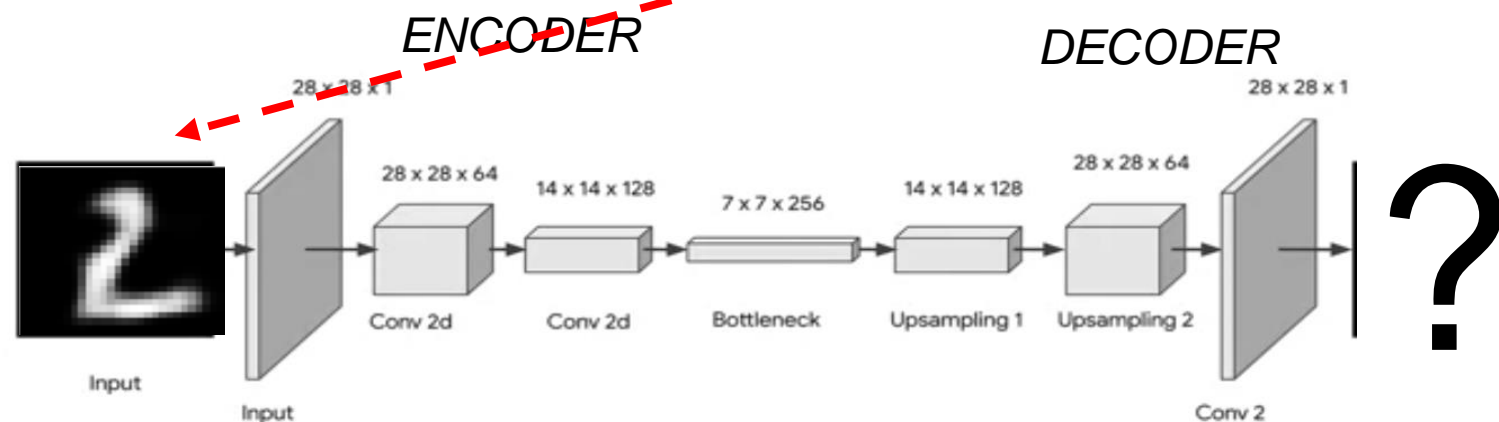
In principle, our denoising autoencoder removed noise pixels and/or filled in digit pixels



In principle, our denoising autoencoder removed noise pixels and/or filled in digit pixels

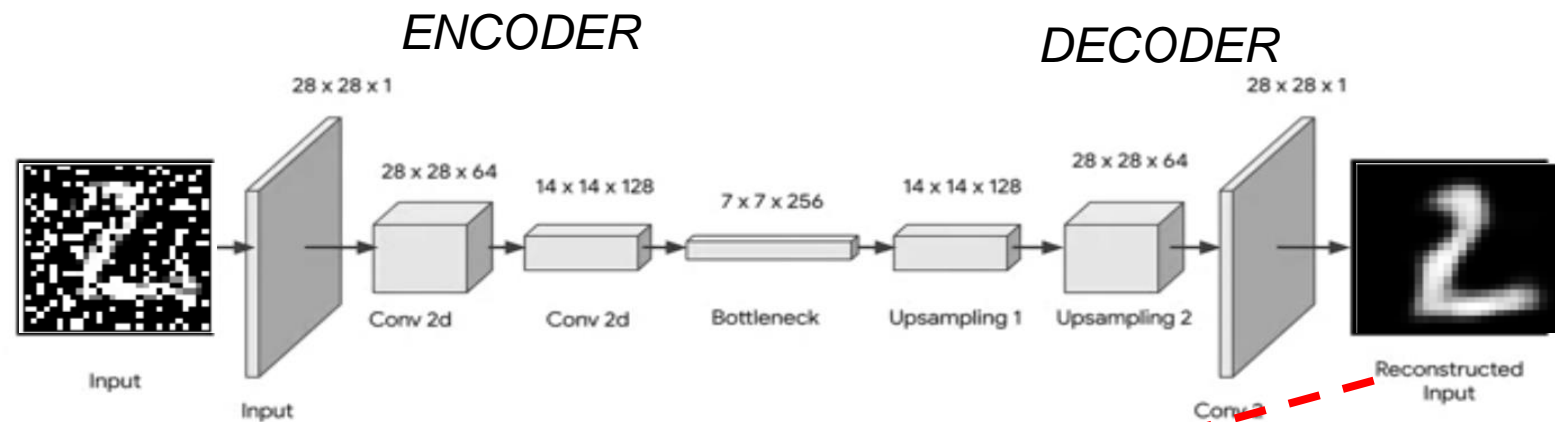


What would happen if we fed the denoised output back into the autoencoder?





In principle, our denoising autoencoder removed noise pixels and/or filled in missing pixels




What would happen if we fed the denoised output back into the autoencoder?


A: better reconstruction

B: all pixels would be removed

C: all pixels would be filled in




First step of denoising



first step of denoising




1 more step of denoising




Is it better?

First step of denoising




3 more steps of denoising



**Frist step of denoising**



**3 more steps of denoising**



**Let's make this more stable, by training a network to just remove a little noise. It is like training to predict noise diffusion.**

# Stable Diffusion for Image Reconstruction

Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. 2020.

- **Concept:**

create a sequence of images with noise,  $t=1\dots T$



# Stable Diffusion for Image Reconstruction

Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. 2020.


- **Concept:**

create a sequence of images with noise,  $t=1\dots T$



train the network to reconstruct image  $t-1$  from image  $t$

Note: this example is in pixel space, but it is often applied in embedding space



- From Ho et al. 2020

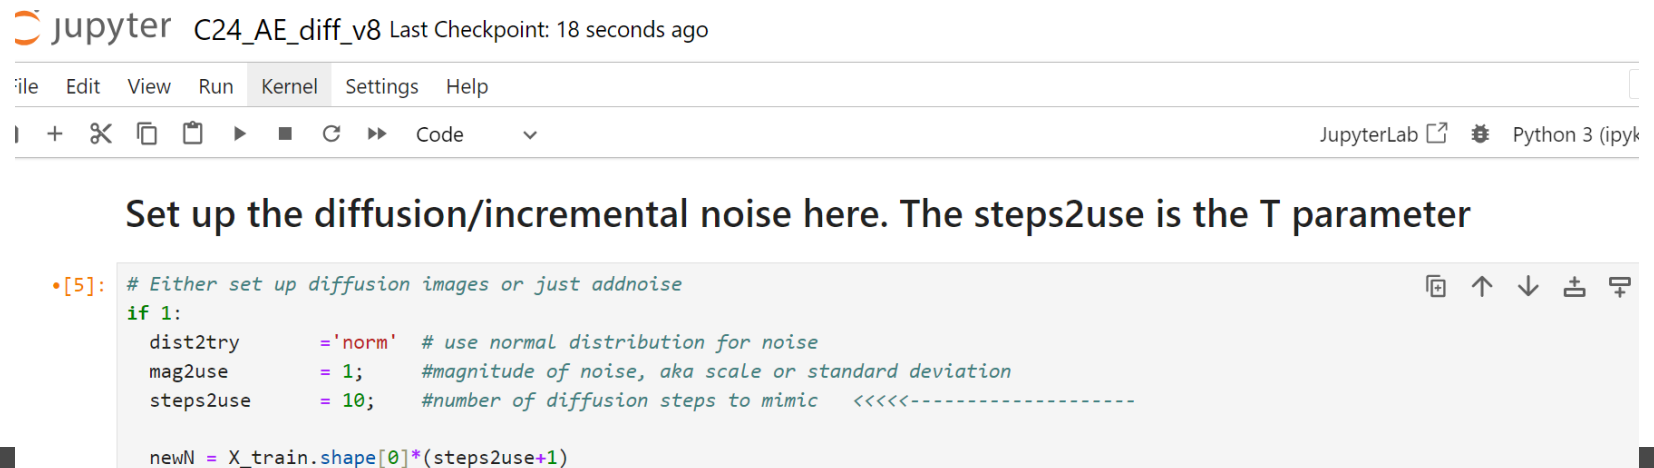
Early denoising steps add overall structure  
Later denoising steps add more detail





# Exercise

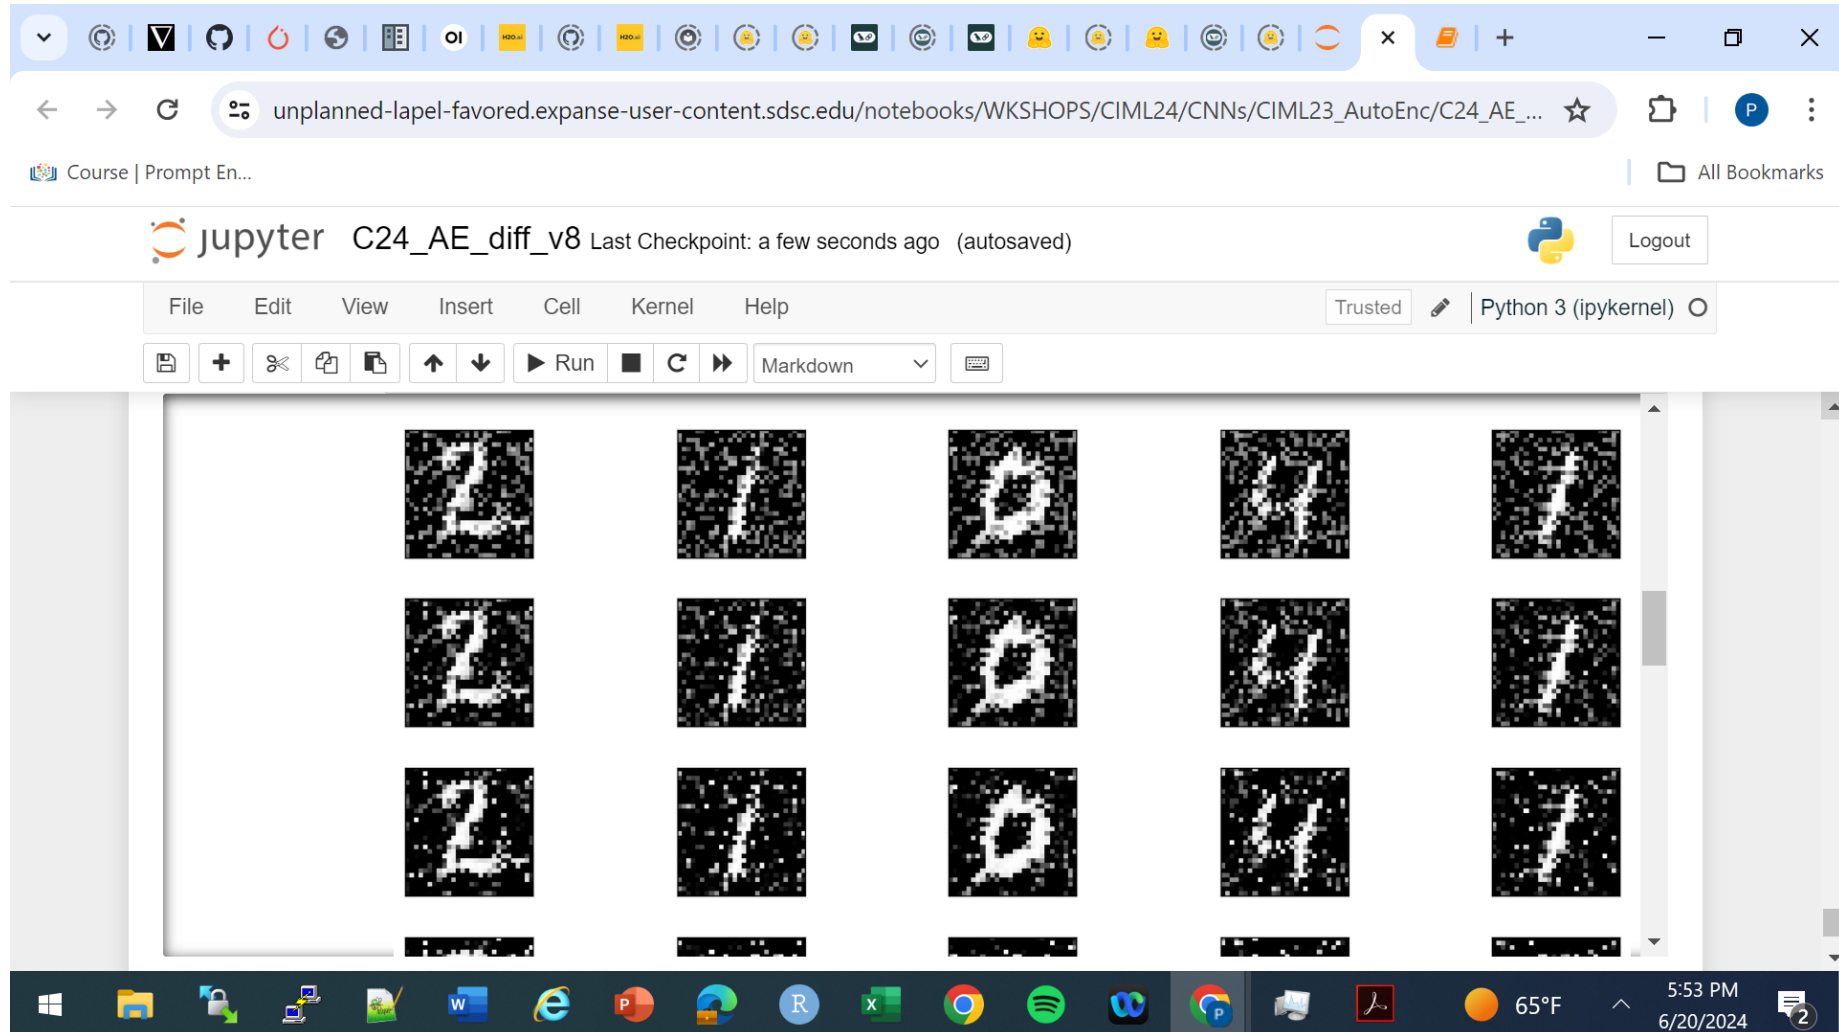
- MNIST stable diffusion, or incremental denoising
- Open and run the notebook
- Try changing T parameter (steps2use)



```
•[5]: # Either set up diffusion images or just addnoise
if 1:
    dist2try      = 'norm' # use normal distribution for noise
    mag2use       = 1;     #magnitude of noise, aka scale or standard deviation
    steps2use     = 10;    #number of diffusion steps to mimic <<<<-----
    newN = X_train.shape[0]*(steps2use+1)
```

Sample output where T1,T2, are going down the columns

What would happen if the input was completely random?



end