

# SDSC HPC/DS Summer Institute 2024

## Deep Learning - Transfer Learning



# Deep Learning Agenda

**8:00 - 8:05 – Welcome**

**8:05 - 9:05 – Intro to NN/CNN**

**9:05 - 9:15 – Break**

**9:15 - 10:15 – Deep Learning**

**10:15 - 11:00 – DL Layers & Architectures**

**11:00 - 11:30 – Break/Lunch**

**11:30 - 12:30 – DL Transfer Learning**

**12:30 - 12:40 – Break**

**12:40 - 1:40 – DL Other Topics**

**1:40 - 2:00 – Wrapup**

# DEEP LEARNING OVERVIEW

- **Neural Network Basics**
  - Processing Unit
  - Activation Function
  - Loss Function
- **Deep Learning Fundamentals**
  - Deep Network Layers
  - DL Architectures
  - DL Libraries
- **Transfer Learning**
  - Transfer Learning Concepts
  - Transfer Learning Demo

# Deep Learning Transfer Learning

Mai H. Nguyen, Ph.D.

# Transfer Learning

- **To overcome challenges of training model from scratch:**
  - Insufficient data
  - Very long training time
- **Use pre-trained model**
  - Trained on another dataset
  - This serves as starting point for model
  - Then train model on current dataset for current task

# Transfer Learning Approaches

- **Feature extraction**

- Remove classification layer from pre-trained model
- Treat rest of network as feature extractor
- Use features to train new classifier
  - “top model” or “classification head”

- **Fine tuning**

- Tune weights in some layers of original model (along with weights of top model)
- Train model for current task using new dataset

# CNNs for Transfer Learning

- **Popular architectures**
  - AlexNet
  - GoogLeNet
  - VGGNet
  - ResNet
- **All winners of ILSVRC**
  - ImageNet Large Scale Visual Recognition Challenge
  - Annual competition on vision tasks on ImageNet data

# ImageNet

- **Database**

- Developed for computer vision research
- ~14,000,000 images hand-annotated
- ~22,000 categories


- **ILSVRC History**

- Started in 2010
- Image classification task: 1,000 object categories
- Image classification error rate
  - 2010: 28.20% (conventional image processing techniques)
  - 2012: 15.30% (AlexNet)
  - 2015: 3.57% (ResNet; better than human performance)
  - 2016: 2.99% (16.7% error reduction)
  - 2017: 2.25% (23.3% error reduction)



# Results on ImageNet Classification

## Classification Results (CLS)




# Transfer Learning

*Input*




*Output*

*Learned  
hierarchy*




*Lee et al. 'Convolutional Deep Belief Networks for Scalable  
Unsupervised Learning of Hierarchical Representations' ICML 2009*

# Pre-Trained Model




<https://www.cs.toronto.edu/~frossard/post/vgg16/>

# Transfer Learning - Feature Extraction



<https://www.cs.toronto.edu/~frossard/post/vgg16/>

# Transfer Learning - Fine Tuning



<https://www.cs.toronto.edu/~frossard/post/vgg16/>

# When & How to Fine Tune

- **New dataset is small & similar to original dataset**
  - Extract features from higher layer and feed to separate classifier
- **New dataset is large & similar to original dataset**
  - Fine tune top or all layers
- **New dataset is small & different from original dataset**
  - Extract features from lower layer and feed to separate classifier
- **New dataset is large & different from original dataset**
  - Fine tune top or all layers

<http://cs231n.github.io/transfer-learning/>

# Practical Tips for Transfer Learning

- **Learning rate**
  - Use very small learning rate for fine tuning. Don't want to destroy what was already learned.
- **Start with properly trained weights**
  - Train top-level classifier first, then fine tune lower layers.
  - Top model with random weights may have negative effects on when fine tuning weights in pre-trained model
- **Data augmentation**
  - Simple ways to slightly alter images
    - Horizontal/vertical flips, random crops, translations, rotations, etc.
  - Use to artificially expand your dataset

# Transfer Learning Hands-On

- **Data**

- Cats and dogs images from Kaggle

- **Exercises**

- Feature extraction
  - Use pre-trained CNN to extract features from images
  - Train neural network to classify cats/dogs using extracted features
  - Code: `feature_extract.ipynb`, `feature_extract_soln.ipynb`
- Fine tune
  - Adjust weights of last few layers of pre-trained CNN and top classifier model through training
  - Code: `finetune.ipynb`, `finetune_soln.ipynb`
- Note
  - Shut down kernel for `feature_extract.ipynb` before running `finetune.ipynb` to avoid out-of-memory errors (Kernel -> Shut Down Kernel)



# Data

- **Subset of Dogs Vs. Cats dataset from Kaggle**
  - <https://www.kaggle.com/c/dogs-vs-cats>
- **Train**
  - 1000 cats + 1000 dogs
- **Validation**
  - 200 cats + 200 dogs
- **Test**
  - 200 cats + 200 dogs



# TRANSFER LEARNING - FEATURE EXTRACTION


- **Data**

- Cats and dogs images from Kaggle

- **Method**

- Use VGG16 trained on ImageNet data as pre-trained model. Remove last fully connected layer.
- Extract features from pre-trained model and save
- Neural network then trained on extracted features to classify cats vs. dogs

# TRANSFER LEARNING - FEATURE EXTRACTION



Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

# Feature Extraction Overview

- **Data**
  - Set image dimensions & location
  - Read images from folder in batches
- **Model**
  - Load model pre-trained on ImageNet data
  - Freeze weights in pre-trained model to use as feature extractor
  - Add top model to classify cats vs dogs
  - Model = Pre-trained base model + top model classifier
- **Train model**
  - Use training data to adjust top model weights
- **Evaluate model**
  - Calculate accuracy, etc.
  - Perform inference on test images

# Setup

- **Login to Expanse**

- Open terminal window on local machine
- `ssh login.expanse.sdsc.edu -l <account>`

- **Pull latest from repo**

- `git pull`
- URL:

<https://github.com/sdsc/sdsc-summer-institute-2024.git>

# Server Setup for TensorFlow - Portal

- **Expanse Portal**
  - <https://portal.expanse.sdsc.edu>
- **Parameters**
  - Account: sds184
  - Time limit (min): 180
  - Number of cores: 10
  - Memory required per node: 93 GB
  - GPUs: 1
  - Singularity image:  
/cm/shared/apps/containers/singularity/ciml/2021/tensorflow-lat  
est.sif
  - Environment module: singularitypro
  - Reservation: ciml-day3
  - Working directory: home
  - Type: JupyterLab

# Server Setup for TensorFlow - Command Line

- **In terminal window**
  - `jupyter-gpu-shared-tensorflow`
    - Alias for:
    - `galileo launch --account ${SI24_ACCOUNT} --reservation ${SI24_RES_GPU} --partition gpu-shared --qos ${SI24_QOS_GPU} --cpus 10 --memory 92 --gpus 1 --time-limit 04:00:00 --env-modules singularitypro --sif ${SI24_CONTAINER_DIR}/tensorflow/tensorflow-latest.sif --bind /cm,/expance,/scratch --nv --quiet`
- **To check queue**
  - `squeue -u $USER`

# Data Setup

- In terminal window in Jupyter Lab, do the following
- Go to your home directory

```
cd
```

```
pwd
```

```
# Should see /home/$USER
```

- Get data

- mkdir data

```
# If doesn't already exist
```

- cd data

- cp /cm/shared/examples/sdsc/ciml/2022/catsVsDogs.zip .

Don't forget the  
period at the end!



- unzip -q catsVsDogs.zip

- ls catsVsDogs

```
# Should see train, val, test
```



# Data

- In terminal window in Jupyter Lab, do the following
- **Get counts of images**
  - `ls`      => Should see data
  - `ls -l data/catsVsDogs/train/cats/* | wc -l`
  - `ls -l data/catsVsDogs/train/dogs/* | wc -l`
  - `ls -l data/catsVsDogs/val/cats/* | wc -l`
  - `ls -l data/catsVsDogs/val/dogs/* | wc -l`
  - `ls -l data/catsVsDogs/test/cats/* | wc -l`
  - `ls -l data/catsVsDogs/test/dogs/* | wc -l`

# TRANSFER LEARNING - FINE TUNING


- **Data**

- Cats and dogs images from Kaggle

- **Method**

- Use VGG16 trained on ImageNet data as pre-trained model.
- Replace last fully connected layer with neural network trained from Feature Extraction hands-on.
- Fine tune last convolution block and fully connected layer.

# TRANSFER LEARNING - FINE TUNING




Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

# Fine Tune Overview


- **Data**
  - Set image dimensions & location
  - Read images from folder in batches
- **Model**
  - Load trained model from feature extraction code
  - Weights in last few convolutional blocks and top model will be adjusted during training
  - All other weights in pre-trained model are frozen
- **Train model**
  - Use training data to adjust top model weights
  - Use validation data to determine when to stop training
- **Evaluate model**
  - Calculate accuracy, etc.
  - Perform inference on test images

# Early Stopping

Using validation data to determine when to stop training to avoid overfitting



# Data Augmentation

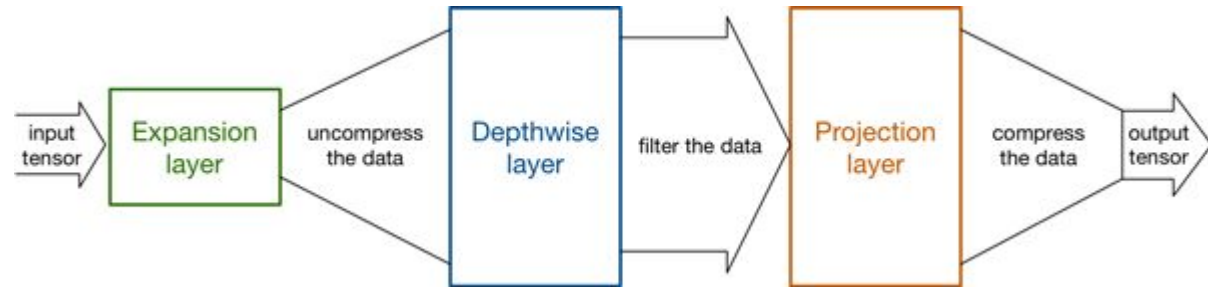
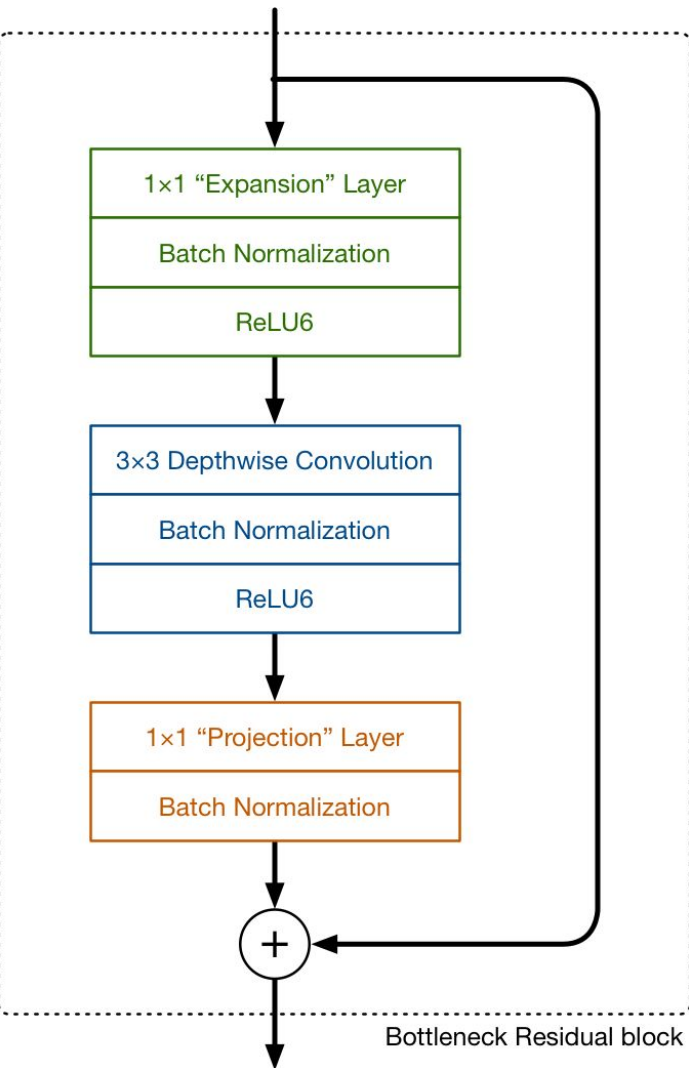


Add variability to your dataset

<https://nanonets.com/blog/data-augmentation-how-to-use-deep-learning-when-you-have-limited-data-part-2/>

# MobileNetV2


- CNN
- Lightweight architecture
- Designed for mobile devices



<https://machinethink.net/blog/mobilenet-v2/>

# MobileNetV2

- CNN
- Lightweight architecture
- Designed for mobile devices





# RESOURCES

- **TensorFlow Tutorial on Transfer Learning**
  - [https://www.tensorflow.org/tutorials/images/transfer\\_learning](https://www.tensorflow.org/tutorials/images/transfer_learning)
- **Transfer Learning**
  - <http://cs231n.github.io/transfer-learning/>
- **ImageNet**
  - <http://www.image-net.org>
- **TensorFlow/Keras API**
  - [https://www.tensorflow.org/api\\_docs/python/tf/keras/Model](https://www.tensorflow.org/api_docs/python/tf/keras/Model)